

Postgres 14 WAL Replication: Easy Step-By-Step Guide



Overview

This blog is a simple walk-through of the WAL Streaming replication using the latest PostgreSQL-14 on Ubuntu 20.04 focal:

- Setups PostgreSQL 14 on primary and secondary.
- Configure primary-secondary.
- Create a replica user on the primary with replication privileges.
- Copy data files from primary to secondary.
- Test.

Why WAL replication!

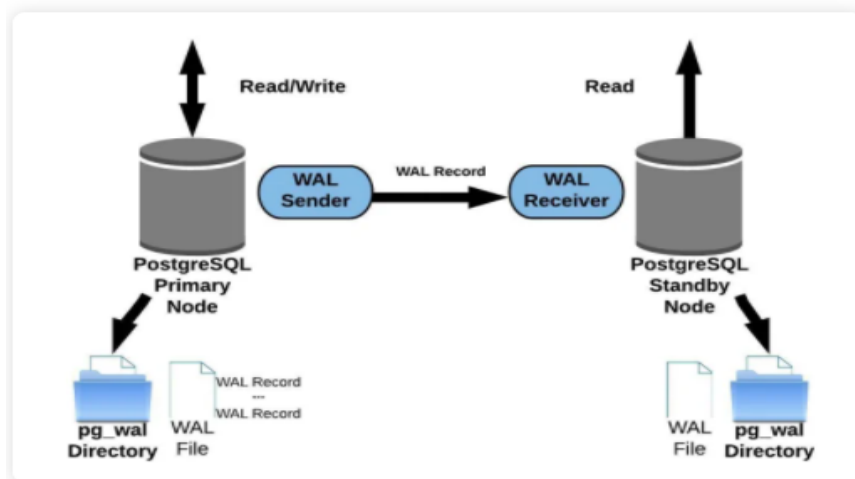
Data replication ensures the duplication of data on an ongoing basis so that replication is in a consistently updated state and is synchronized with the source. This ensures high data availability and safety of data against undesired events such as crashes, system errors, etc.

Postgres WAL replication meets this requirement through a feature called streaming replication. This feature is achieved through a master-slave configuration.

Prerequisite

- 2 Ubuntu server version 20.04.3 LTS.
- Port 5432 is enabled between servers.
- PostgreSQL 14 running on both.

Streaming WAL Records



Postgres 14 WAL Replication: Easy Step-By-Step Guide



➤ Install Postgresql 14 on primary and secondary:

- **Note the below steps should be applied on both.**

- Create the file repository configuration:

- Import the repository signing key:

```
root@primary-server:/var/lib# sudo sh -c 'echo "deb http://apt.postgresql.org/pub/
repos/apt $(lsb_release -cs)-pgdg main" > /etc/apt/sources.list.d/pgdg.list'
root@primary-server:/var/lib# wget --quiet -O - https://www.postgresql.org/media/k
eys/ACCC4CF8.asc | sudo apt-key add -
OK
```

- Update the package lists:

```
root@primary-server:/var/lib# sudo apt-get update
Hit:1 http://archive.ubuntu.com/ubuntu focal InRelease
Hit:2 http://archive.ubuntu.com/ubuntu focal-updates InRelease
Hit:3 http://apt.postgresql.org/pub/repos/apt focal-pgdg InRelease
Hit:4 http://archive.ubuntu.com/ubuntu focal-backports InRelease
Hit:5 http://archive.ubuntu.com/ubuntu focal-security InRelease
Hit:6 https://repos.citusdata.com/community/ubuntu focal InRelease
Reading package lists... Done
```

- Install the latest version of PostgreSQL. If you want a specific version, use 'postgresql-12' or similar instead of 'postgresql':

```
root@primary-server:/var/lib# sudo apt-get -y install postgresql
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following additional packages will be installed:
  pgdg-keyring postgresql-14 postgresql-client-14 postgresql-client-common
  postgresql-common
```

- **Note: make sure you issue the below command which makes sure PostgreSQL starts automatically when server get rebooted.**

```
root@primary-server:/var/lib# systemctl enable postgresql
Synchronizing state of postgresql.service with SysV service script with /lib/syste
md/systemd-sysv-install.
Executing: /lib/systemd/systemd-sysv-install enable postgresql
```

- Let's check Postgresql status:

```
root@primary-server:/var/lib# systemctl status postgresql
● postgresql.service - PostgreSQL RDBMS
   Loaded: loaded (/lib/systemd/system/postgresql.service; enabled; vendor pres
   Active: active (loaded) since Mon 2021-10-25 20:54:25 UTC; 8min ago
   Main PID: 403478 (code=exited, status=0/SUCCESS)
     Tasks: 0 (limit: 19108)
    Memory: 0B
     CGroup: /system.slice/postgresql.service

Oct 25 20:54:25 primary-server systemd[1]: Starting PostgreSQL RDBMS...
Oct 25 20:54:25 primary-server systemd[1]: Finished PostgreSQL RDBMS.
```

➤ Configure Primary server:

Postgres 14 WAL Replication: Easy Step-By-Step Guide



- Edit Postgresql.conf and apply the below changes:
- vim /etc/postgresql/main/14/postgresql.conf:
listen_addresses = '*' # what IP address(es) to listen on;
wal_level = replica # minimal, replica, or logical
hot_standby = on # "off" disallows queries during recovery

- Create a user with replication privileges:

```
root@primary-server:/# systemctl restart postgresql
root@primary-server:/# sudo -u postgres psql
psql (14.0 (Ubuntu 14.0-1.pgdg20.04+1))
Type "help" for help.

postgres=# CREATE USER replica REPLICATION LOGIN ENCRYPTED PASSWORD 'replica';
CREATE ROLE
```

- Then we need to add the permission to **pg_hba.conf** file to enable a Standby server to access a Primary server.
- vim /etc/postgresql/14/main/pg_hba.conf:
host replication replica SecondaryIP/32
md5
- Make sure you restart Postgresql after applying these changes.

➤ Configure Secondary server:

- Edit Postgresql.conf and apply the below changes:
- vim /etc/postgresql/main/14/postgresql.conf:
listen_addresses = '*' # what IP address(es) to listen on.

- Remove data directory from secondary server

```
root@secondary-server:/home/ubuntu# sudo su -l postgres
postgres@secondary-server:~$ rm -rfv /var/lib/postgresql/14/main/*
removed '/var/lib/postgresql/14/main/base/1/1255'
```

- Let's take backup from the primary data directory as shown below:

```
postgres@secondary-server:~/14/main$ pg_basebackup -h 192.168.1.100 -p 5432 -U replica -D /var/lib/postgresql/14/main/ -Fp -Xs -R
```

- Password: █

Notes:

- It promotes you to enter the replica password created on the primary server.
- -R creates the replication file configuration automatically as shown below:

Postgres 14 WAL Replication: Easy Step-By-Step Guide



```
root@secondary-server:/var/lib/postgresql/14/main# cat postgresql.auto.conf
# Do not edit this file manually!
# It will be overwritten by the ALTER SYSTEM command.
primary_conninfo = 'user=replica password=replica password channel_binding=prefer host=192.168.1.100 port=5432 sslmode=prefer sslcompression=0 sslsni=1 ssl_min_protocol_version=TLSv1.2 gssencmode=prefer krbsrvname=postgres target_session_attrs=any'
```

- Make sure you restart Postgresql and check its status:

```
root@secondary-server:/home/ubuntu# systemctl restart postgresql
root@secondary-server:/home/ubuntu# systemctl status postgresql
● postgresql.service - PostgreSQL RDBMS
   Loaded: loaded (/lib/systemd/system/postgresql.service; enabled; vendor preset:
   Active: active (exited) since Mon 2021-10-25 21:49:35 UTC; 6s ago
   Process: 389964 ExecStart=/bin/true (code=exited, status=0/SUCCESS)
   Main PID: 389964 (code=exited, status=0/SUCCESS)

Oct 25 21:49:35 secondary-server systemd[1]: Starting PostgreSQL RDBMS...
Oct 25 21:49:35 secondary-server systemd[1]: Finished PostgreSQL RDBMS.
```

- Now let's create database and keep an eye on the secondary server:

Primary

```
postgres=# create database hello_world;
CREATE DATABASE
postgres=# \l
```

Name	Owner	Encoding	Collate	Ctype	Access privileges
hello_world	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
template0	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	=c/postgres postgres=CTc/postgres
template1	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	=c/postgres postgres=CTc/postgres

Secondary

```
postgres=# \l
```

Name	Owner	Encoding	Collate	Ctype	Access privileges
hello_world	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
postgres	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
template0	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	=c/postgres postgres=CTc/postgres
template1	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	=c/postgres postgres=CTc/postgres

- Let's import a dump backup into primary to test the replication is streaming as expected:

```
postgres=# create database data;
CREATE DATABASE
postgres=# \q
postgres@primary-server:/home/ubuntu$ gunzip -c data_export_11_8_2021.sql.gz | psql data
```

- Check the secondary if the database exists:

```
postgres=# \l
```

Name	Owner	Encoding	Collate	Ctype	Access privileges
data	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	
hello_world	postgres	UTF8	en_US.UTF-8	en_US.UTF-8	

- Let's check the size of the database in both:

Primary

```
postgres=# SELECT pg_size_pretty(pg_database_size('data'));
pg_size_pretty
-----
1419 MB
(1 row)
```

Secondary

```
postgres=# SELECT pg_size_pretty(pg_database_size('data'));
pg_size_pretty
-----
1419 MB
(1 row)
```

Postgres 14 WAL Replication: Easy Step-By-Step Guide



- You can also check on the primary to find out the state of streaming:

- `select * from pg_stat_replication;`

pid	usesysid	username	application_name	client_addr	client_hostname	client_port	backend_start	backend_xmin	state
sent_lsn	write_lsn	flush_lsn	replay_lsn	write_lag	flush_lag	replay_lag	sync_priority	sync_state	reply_time
408892	16384	replica	14/main			52944	2021-10-25 21:49:33.068196+00		streaming
/A8C31FC0	0/A8C31FC0	0/A8C31FC0	0/A8C31FC0				0	async	2021-10-25 22:23:41.891326+00
(1 row)									

- On secondary server you check by:

- `select * from pg_stat_wal_receiver;`

pid	status	receive_start_lsn	receive_start_tli	written_lsn	flushed_lsn	received_tli	last_msg_send_time	last_msg_receipt_time
latest_end_lsn	latest_end_time	slot_name	sender_host	sender_port	conninfo			
389951	streaming	0/3000000	1	0/A8C31FC0	0/A8C31FC0	1	2021-10-25 22:28:22.30422+00	2021-10-25 22:28:22.405516+00
0/A8C31FC0	2021-10-25 22:14:51.050281+00			5432	user=replica password=***** channel_binding=prefer dbname=replicatio			
n host=	port=5432 fallback_application_name=14/main sslmode=prefer sslcompression=0 sslsn=1 ssl_min_protocol_version=TLSv1.2 gssencmode=prefer krbsrvna							
me=postgres target_session_attrs=any								
(1 row)								

Now you have a primary server and streaming replication performing as it should be.

Hope this helps!!

Postgres 14 WAL Replication: Easy Step-By-Step Guide



➤ References:

1. <https://hevodata.com/learn/postgres-wal-replication/#master>
2. <https://www.postgresql.org/download/linux/ubuntu/>